




Hyperion Steward

Challenge 1 · HYPER-AI | says done only once it has seen it done



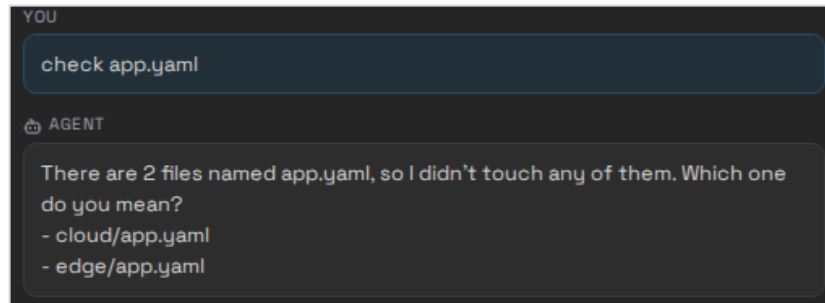
Funded by
the European Union

GitHub repo

- **GitHub repo:** github.com/Faadil1/veles-hack-2026
- **Docker image:** `faadil12/hyperion:latest` (exposes :8000/chat)
- **Run:** `docker run -p 8000:8000 --add-host host.docker.internal:host-gateway -e API_KEY=<key> faadil12/hyperion:latest`
- **Licence:** Apache-2.0, as in the official hyperion-starter
- **Evidence:** CI reports on the ci-evidence branch (official IDE images, digests recorded)

Summary

- A drop-in /chat microservice for the HYPER-AI IDE: answers HYPER-AI questions from the official documents and turns plain language into IDE actions.
- We read the IDE the judges run, not only its tutorial. Its actions are fire-and-forget: a naive agent says done when nothing happened.
- Steward checks every profile with an exact copy of the IDE validator, then reads the workspace back and only claims what it saw. Deletes always ask; overwrites ask unless you named the file or Steward created it; file changes can be undone (a folder delete cannot).



The official IDE (organisers' GUI and backend images), a browser typing into the Hyperion panel in CI.

Highlights

0

disagreements with the IDE's own validator, over 5,536 test profiles

9/9

safety scenarios handled correctly (same tool calls, stub mirroring the IDE). Naive agent: 1/9

0

false "done" claims in the same scenarios. Naive agent: 5

6/6

live steps against the official backend image, every change verified by read-back

Also: works without a model for the official example requests (create, delete, check, undo); plain-text replies because the panel does not render markdown; per-session memory sized to the provided 8k-token model.